

Inflation Money

2026-05-15

An Introduction

I've finally came around to checking out Polymarket. To those that don't know it, it's a simple betting site. I wanted to try and learn about time-series analysis and trying a bet on inflation seemed like an easy entry¹

The Bet

The site poses a simple question:

How high will inflation get in 2026?

With the following odds:

\$945,213 Vol. Dec 31, 2026		Polymarket	
Above 4% \$228,871 Vol.	97% ▲ 63%	Buy Yes 97.1¢	Buy No 3.0¢
Above 4.5% \$8,642 Vol.	61% ▲ 11%	Buy Yes 61¢	Buy No 40¢
Above 5% \$137,773 Vol.	32% ▲ 12%	Buy Yes 32¢	Buy No 69¢
Above 6% \$38,218 Vol.	17% ▼ 17%	Buy Yes 18¢	Buy No 84¢
Above 10% \$37,856 Vol.	6%	Buy Yes 7.9¢	Buy No 95.2¢
Above 8% \$37,109 Vol.	7%	Buy Yes 8.4¢	Buy No 93.6¢
Hide resolved ^			
Above 3% \$280,045 Vol.		Yes ✓	
Above 3.5% \$176,699 Vol.		Yes ✓	

We can see that inflation already crossed above 3% and 3.5% - they appear at the bottom as “resolved”, i.e whoever bet on these won. It's important to understand the context of the bet - if at ANY month during the year inflation passes some threshold, it will be resolved in favor of whoever placed a bet on that threshold. We need to get a little formal so we can start thinking about the problem:

¹in hindsight it's probably one of the toughest things to try and predict. The US government tries and fails at it. Clearly, I'll manage to outsmart 'em.

Denote π_m the Consumer Price Index (denoted by CPI from now on) at the end of month $m \in (1, \dots, 12)$. Denote $t \in (4, 4.5, 5, 6, 8, 10) := T$ the critical thresholds suggested in the bet. Our event of interest is defined as $E_t := \{\max_{m \in (1, \dots, 12)} \pi_m\} > t$ for all thresholds t . We would like to model the probability $P(E_t)$ for all $t \in T$

The Data

After understanding what we would like to model, we can start building some basic models. First, we need to gather some data. I've downloaded data about past CPI information². Each row is a single year, each column a month. As always, we need to clean up the data before we can start using it effectively. The following code imports the data, and cleans it up so we can start processing it:

```
data <- read_excel("hist_cpi-u.xlsx", skip = 3)
data <- data[-1, ] # remove first row
data <- data[, -1] # remove first col
data <- data %>%
  slice(1:(n()-3)) %>% # remove last 3 lines
  mutate(across(everything(), as.numeric)) # make sure all columns are numeric
```

here are the first couple of rows to get a feel for what we have:

```
## # A tibble: 6 x 13
##   Year Jan. Feb. Mar. Apr. May Jun. Jul. Aug. Sep. Oct. Nov. Dec.
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1 1913  9.8  9.8  9.8  9.8  9.7  9.8  9.9  9.9  10  10  10.1  10
## 2 1914  10  9.9  9.9  9.8  9.9  9.9  10  10.2  10.2  10.1  10.2  10.1
## 3 1915  10.1 10  9.9  10  10.1  10.1  10.1  10.1  10.1  10.2  10.3  10.3
## 4 1916  10.4 10.4 10.5 10.6 10.7 10.8 10.8 10.9 11.1 11.3 11.5 11.6
## 5 1917  11.7 12  12  12.6 12.8 13  12.8 13  13.3 13.5 13.5 13.7
## 6 1918  14  14.1 14  14.2 14.5 14.7 15.1 15.4 15.7 16  16.3 16.5
```

The Models

Now let's get our hands dirty with time-series analysis. In general, a time-series is just a sequence of data points indexed in chronological order. We can order our data in many different ways. One natural way to do so, is to go linearly in time. We can take $x_0 := PCI_{\text{January, 1913}}$, $x_1 := PCI_{\text{February, 1913}}$ and so on.

²<https://www.bls.gov/cpi/tables/supplemental-files/>